

# Homework 3

[Ra-Zakee Muhammad]

Due 9/21/2021

**Classmates/other resources consulted:** [type answer here]

```
library(ggplot2)
library(dplyr)
library(nycflights13)
```

## Question 1 (3 points)

Explain why the following results return FALSE, and how you should compare these values instead. You can give one explanation for both, you do not need to give a separate explanation for each. (note: every computer is different, and while they both return FALSE on my computer, they may not both return FALSE on your computer. Regardless, discuss why FALSE might show up as an output and what you should do instead).

```
?near()
```

```
sqrt(3)^2 == 3
```

```
## [1] FALSE
```

```
1.45 - 0.55 == 2.45 - 1.55
```

```
## [1] FALSE
```

These are different data types, one is an integer another is a dbl therefore in order to compare the values one needs to use the near function.

```
near(sqrt(3)^2, 3)
```

```
## [1] TRUE
```

```
near(1.45 - 0.55 , 2.45 - 1.55)
```

```
## [1] TRUE
```

## Question 2 (12 points)

Using the flights data set, output a tibble consisting of flights meeting the following criteria. For each, indicate how many flights there are.

- a. flights with an arrival delay of more than two hours

```
flights %>% filter(!is.na(arr_delay), arr_delay > 2)
```

```
## # A tibble: 123,096 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     554           558          -4     740           728
## 5  2013     1     1     555           600          -5     913           854
## 6  2013     1     1     558           600          -2     753           745
## 7  2013     1     1     558           600          -2     924           917
## 8  2013     1     1     559           600          -1     941           910
## 9  2013     1     1     600           600           0     837           825
## 10 2013     1     1     602           605          -3     821           805
## # ... with 123,086 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

- b. flights operated by United (UA), American (AA), or Delta (DL) (whether or not they took off)

```
flights %>%
  filter( carrier == "UA" | carrier == "AA" | carrier == "DL" )
```

```
## # A tibble: 139,504 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     554           600          -6     812           837
## 5  2013     1     1     554           558          -4     740           728
## 6  2013     1     1     558           600          -2     753           745
## 7  2013     1     1     558           600          -2     924           917
## 8  2013     1     1     558           600          -2     923           937
## 9  2013     1     1     559           600          -1     941           910
## 10 2013     1     1     559           600          -1     854           902
## # ... with 139,494 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

- c. flights *not* operated by United, American, or Delta (whether or not they took off)

```
flights %>%
  filter( carrier != "UA" , carrier != "AA" , carrier != "DL" )
```

```
## # A tibble: 197,272 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     544           545          -1    1004           1022
## 2  2013     1     1     555           600          -5     913           854
## 3  2013     1     1     557           600          -3     709           723
## 4  2013     1     1     557           600          -3     838           846
## 5  2013     1     1     558           600          -2     849           851
## 6  2013     1     1     558           600          -2     853           856
## 7  2013     1     1     559           559           0     702           706
## 8  2013     1     1     600           600           0     851           858
## 9  2013     1     1     600           600           0     837           825
## 10 2013     1     1     601           600           1     844           850
## # ... with 197,262 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

d. flights that were scheduled to depart in February, March, May, July, September, or November (use %in%)

```
flights %>%
  filter( month %in% c(2,3,5,7,9,11))
```

```
## # A tibble: 166,848 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013    11     1         5           2359           6     352           345
## 2  2013    11     1        35           2250        105     123           2356
## 3  2013    11     1      455           500          -5     641           651
## 4  2013    11     1      539           545          -6     856           827
## 5  2013    11     1      542           545          -3     831           855
## 6  2013    11     1      549           600         -11     912           923
## 7  2013    11     1      550           600         -10     705           659
## 8  2013    11     1      554           600          -6     659           701
## 9  2013    11     1      554           600          -6     826           827
## 10 2013    11     1      554           600          -6     749           751
## # ... with 166,838 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

e. flights that departed between 7am and 8am, including both 7am and 8am

```
flights %>%
  filter(dep_time >= 700 , dep_time <= 800) %>% arrange((dep_time))
```

```
## # A tibble: 22,009 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
```

```
## 1 2013 1 2 700 630 30 917 840
## 2 2013 1 2 700 700 0 851 850
## 3 2013 1 2 700 700 0 1017 1015
## 4 2013 1 3 700 700 0 851 836
## 5 2013 1 4 700 700 0 941 1025
## 6 2013 1 4 700 700 0 956 1025
## 7 2013 1 5 700 700 0 1014 1025
## 8 2013 1 6 700 615 45 1001 921
## 9 2013 1 7 700 704 -4 912 932
## 10 2013 1 7 700 700 0 959 1045
## # ... with 21,999 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

### Question 3 (6 points)

a. Why does `TRUE | NA` not result in `NA`?

because the command `|` only requires that one of the values on either side be true to return a true value

b. Why does `NA & FALSE` not result in `NA`?

Because when using the “&” or “,” command if either side has a false value then the command returns false

c. Why does `NA^0` not result in `NA`?

Probably because the r language reads any value numeric or logical raised to the 0th power as 1.

```
FALSE^0
```

```
## [1] 1
```

```
TRUE^0
```

```
## [1] 1
```

### Question 4 (6 points)

Use `arrange()` to do the following

a. Sort flights to find the flights with the shortest air time. Among the ten flights with the shortest airtime, where did they go?

```
flights %>%
  filter(!is.na(air_time)) %>%
  arrange(air_time) %>%
  select(air_time, dest) %>%
  head(10) %>%
  mutate(n = row_number())
```

```
## # A tibble: 10 x 3
##   air_time dest      n
##   <dbl> <chr> <int>
## 1      20 BDL      1
## 2      20 BDL      2
## 3      21 BDL      3
## 4      21 PHL      4
## 5      21 BDL      5
## 6      21 PHL      6
## 7      21 BOS      7
## 8      21 PHL      8
## 9      21 BDL      9
## 10     21 BDL     10
```

“Bradley International Airport in Hartford County, Connecticut” ; “Philadelphia International Airport” ;  
 “Boston Logan International Airport”

- b. Sort flights to find the flights that had the longest arrival delay. Among the five flights that had the longest arrival delays, what airport did four of them originate from?

```
flights %>%
  filter(!is.na(arr_delay)) %>%
  arrange(desc(arr_delay)) %>%
  select(arr_delay, origin)%>%
  head(10) %>%
  mutate(n = row_number())
```

```
## # A tibble: 10 x 3
##   arr_delay origin      n
##   <dbl> <chr> <int>
## 1    1272 JFK      1
## 2    1127 JFK      2
## 3    1109 EWR      3
## 4    1007 JFK      4
## 5     989 JFK      5
## 6     931 JFK      6
## 7     915 LGA      7
## 8     895 LGA      8
## 9     878 EWR      9
## 10    875 EWR     10
```

“John F. Kennedy International Airport in Queens, New York” ; “Newark Liberty International Airport”,  
 “Drummond Twins International Airport in Cochrane, Chile”

### Question 5 (9 points)

- a. In the flights data set, give a command that returns all columns/factors, in order, from carrier to destination.

```
flights %>%
  select(carrier: dest)
```

```
## # A tibble: 336,776 x 5
##   carrier flight tailnum origin dest
##   <chr>    <int> <chr>   <chr> <chr>
## 1 UA        1545 N14228  EWR   IAH
## 2 UA        1714 N24211  LGA   IAH
## 3 AA        1141 N619AA  JFK   MIA
## 4 B6         725 N804JB  JFK   BQN
## 5 DL         461 N668DN  LGA   ATL
## 6 UA        1696 N39463  EWR   ORD
## 7 B6         507 N516JB  EWR   FLL
## 8 EV        5708 N829AS  LGA   IAD
## 9 B6          79 N593JB  JFK   MCO
## 10 AA        301 N3ALAA  LGA   ORD
## # ... with 336,766 more rows
```

- b. In the flights data set, give a command that uses “contains” to return the three columns/factors related to arrival time.

```
flights %>%
  select(contains("arr_"))
```

```
## # A tibble: 336,776 x 3
##   arr_time sched_arr_time arr_delay
##   <int>         <int>         <dbl>
## 1     830             819           11
## 2     850             830           20
## 3     923             850           33
## 4    1004            1022          -18
## 5     812             837          -25
## 6     740             728           12
## 7     913             854           19
## 8     709             723          -14
## 9     838             846           -8
## 10    753             745            8
## # ... with 336,766 more rows
```

- c. In the flights data set, write a command that returns all columns/factors except for tailnum, hour, minute, and time\_hour.

```
flights %>%
  select(-tailnum , -hour , -minute, -time_hour)
```

```
## # A tibble: 336,776 x 15
```

```
##      year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##      <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
##  1  2013     1     1     517           515         2     830           819
##  2  2013     1     1     533           529         4     850           830
##  3  2013     1     1     542           540         2     923           850
##  4  2013     1     1     544           545        -1    1004          1022
##  5  2013     1     1     554           600        -6     812           837
##  6  2013     1     1     554           558        -4     740           728
##  7  2013     1     1     555           600        -5     913           854
##  8  2013     1     1     557           600        -3     709           723
##  9  2013     1     1     557           600        -3     838           846
## 10  2013     1     1     558           600        -2     753           745
## # ... with 336,766 more rows, and 7 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>
```

## Question 6 (6 points)

- a. In the mpg data set, add a column/factor that gives the average of the city miles per gallon and the highway miles per gallon

```
mpg %>%
  mutate(average_cty_hwy = (cty + hwy)/2)
```

```
## # A tibble: 234 x 12
##   manufacturer model   displ  year  cyl trans drv   cty   hwy fl   class
##   <chr>         <chr>   <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
## 1 audi         a4       1.8  1999   4 auto~ f     18    29 p   comp~
## 2 audi         a4       1.8  1999   4 manu~ f     21    29 p   comp~
## 3 audi         a4       2    2008   4 manu~ f     20    31 p   comp~
## 4 audi         a4       2    2008   4 auto~ f     21    30 p   comp~
## 5 audi         a4       2.8  1999   6 auto~ f     16    26 p   comp~
## 6 audi         a4       2.8  1999   6 manu~ f     18    26 p   comp~
## 7 audi         a4       3.1  2008   6 auto~ f     18    27 p   comp~
## 8 audi         a4 quattro 1.8  1999   4 manu~ 4     18    26 p   comp~
## 9 audi         a4 quattro 1.8  1999   4 auto~ 4     16    25 p   comp~
## 10 audi        a4 quattro 2    2008   4 manu~ 4     20    28 p   comp~
## # ... with 224 more rows, and 1 more variable: average_cty_hwy <dbl>
```

- b. In the flights data set, add a column/factor that is the larger of the departure delay and the arrival delay.

```
flights %>%
  mutate(max_departure_arrival_delay = pmax(dep_delay, arr_delay))
```

```
## # A tibble: 336,776 x 20
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1     1     517           515         2     830           819
```

```
## 2 2013 1 1 533 529 4 850 830
## 3 2013 1 1 542 540 2 923 850
## 4 2013 1 1 544 545 -1 1004 1022
## 5 2013 1 1 554 600 -6 812 837
## 6 2013 1 1 554 558 -4 740 728
## 7 2013 1 1 555 600 -5 913 854
## 8 2013 1 1 557 600 -3 709 723
## 9 2013 1 1 557 600 -3 838 846
## 10 2013 1 1 558 600 -2 753 745
## # ... with 336,766 more rows, and 12 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   max_departure_arrival_delay <dbl>
```

### Question 7 (12 points)

In this question, you'll need to determine the correct transformation(s) to do on the data set `flights` to produce the desired tibbles and answer the question. *In each part, use pipes.*

- Give a command that returns all the data in `flights` for factors that contain the letter "d" in the factor name, and also includes the month and day of each flight. What are the dimensions of your new tibble?

```
flights %>%
  select(contains("d") | month)
```

```
## # A tibble: 336,776 x 9
##   day dep_time sched_dep_time dep_delay sched_arr_time arr_delay dest
##   <int> <int> <int> <dbl> <int> <dbl> <chr>
## 1 1 517 515 2 819 11 IAH
## 2 1 533 529 4 830 20 IAH
## 3 1 542 540 2 850 33 MIA
## 4 1 544 545 -1 1022 -18 BQN
## 5 1 554 600 -6 837 -25 ATL
## 6 1 554 558 -4 728 12 ORD
## 7 1 555 600 -5 854 19 FLL
## 8 1 557 600 -3 723 -14 IAD
## 9 1 557 600 -3 846 -8 MCO
## 10 1 558 600 -2 745 8 ORD
## # ... with 336,766 more rows, and 2 more variables: distance <dbl>, month <int>
```

- Recall speed in miles per hour is `distance / air_time * 60`. What two airlines had all six of the fastest flights?

```
flights %>%
  mutate(speed = distance / air_time * 60) %>%
  arrange(desc(speed))
```



```
## # A tibble: 336,776 x 20
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     5    25    1709           1700          9    1923           1937
## 2  2013     7     2    1558           1513         45    1745           1719
## 3  2013     5    13    2040           2025         15    2225           2226
## 4  2013     3    23    1914           1910          4    2045           2043
## 5  2013     1    12    1559           1600         -1    1849           1917
## 6  2013    11    17     650            655         -5    1059           1150
## 7  2013     2    21    2355           2358         -3     412            438
## 8  2013    11    17     759            800         -1    1212           1255
## 9  2013    11    16    2003           1925         38      17             36
## 10 2013    11    16    2349           2359        -10     402            440
## # ... with 336,766 more rows, and 12 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   speed <dbl>
```

DL (delta) and EV(express jet)

- c. Consider the flights that arrived more than an hour late but didn't leave late. How many of these were there in September?

```
flights %>%
  filter(arr_delay > 60, dep_delay <= 0, month == 9) %>%
  mutate(n = row_number()) %>%
  summarise(max(n))
```

```
## # A tibble: 1 x 1
##   'max(n)'
##   <int>
## 1      52
```

- d. The information about flights suggests that it should always be the case that  $\text{sched\_dep\_time} + \text{dep\_delay} = \text{dep\_time}$ . Create a new tibble with a column whose value is  $\text{sched\_dep\_time} + \text{dep\_delay}$  and see for how many rows this is different from  $\text{dep\_time}$ . Hypothesize about what might be occurring here. (Hint: use multiple transformations)

```
flights %>%
  mutate(dep_time_1 = sched_dep_time + dep_delay) %>%
  mutate(diff_dep_time = dep_time_1 - dep_time) %>%
  filter(diff_dep_time != 0) %>%
  mutate(n = row_number())
```

```
## # A tibble: 99,777 x 22
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     554            600         -6     812            837
## 2  2013     1     1     555            600         -5     913            854
## 3  2013     1     1     557            600         -3     709            723
## 4  2013     1     1     557            600         -3     838            846
```

```
## 5 2013 1 1 558 600 -2 753 745
## 6 2013 1 1 558 600 -2 849 851
## 7 2013 1 1 558 600 -2 853 856
## 8 2013 1 1 558 600 -2 924 917
## 9 2013 1 1 558 600 -2 923 937
## 10 2013 1 1 559 600 -1 941 910
## # ... with 99,767 more rows, and 14 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   dep_time_1 <dbl>, diff_dep_time <dbl>, n <int>
```

The problem is that R isn't automatically thinking about addition and subtraction in terms of mod 60 (which is the addition that is necessary to maintain the hour structure) while it is helpful that scheduled departure time is in military time, the data type of the column is still integers which means that if we are representing 6:00 as 600 then adding a departure delay of -6 minutes will give us 594 as opposed to 5:54. This explains why many of the non-zero values for `dep_time_1` are 40 because  $50 + 40$  is 90.

### Question 8 (3 points)

Explain why the comparison `x == y` in the following code doesn't produce FALSE, since x and y are different vectors.

```
x <- c(5,2,9,4)
y <- c(5,2,11,6)
x == y
```

```
## [1] TRUE TRUE FALSE FALSE
```

Because in R when you compare vectors using equality or inequalities, the language compares by common indexed entries as opposed to the entire vectors. This way, you get dimension number of logical values if in fact the two vectors share dimension.

### Question 9 (3 points)

**Summarize the starwars dataset: find the minimum mass, maximum height, and median birth year of all characters in the data set (for whom these values are defined). Your answer should be a tibble that only contains these three values.**

```
starwars %>%
  arrange(mass)
```

```
## # A tibble: 87 x 14
##   name      height mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Ratts T~    79    15 none      grey, blue unknown      NA male masculi
## 2 Yoda        66    17 white      green      brown      896 male masculi
## 3 Wicket ~    88    20 brown      brown      brown        8 male masculi
```

```
## 4 R2-D2      96    32 <NA>      white, bl~ red      33 none mascu~
## 5 R5-D4      97    32 <NA>      white, red red      NA none mascu~
## 6 Sebulba   112    40 none      grey, red orange    NA male mascu~
## 7 Dud Bolt   94    45 none      blue, grey yellow    NA male mascu~
## 8 Padmé A~  165    45 brown     light    brown      46 fema~ femin~
## 9 Wat Tam~  193    48 none      green, gr~ unknown    NA male mascu~
## 10 Sly Moo~ 178    48 none      pale     white      NA <NA> <NA>
## # ... with 77 more rows, and 5 more variables: homeworld <chr>, species <chr>,
## #   films <list>, vehicles <list>, starships <list>
```

```
starwars %>%
  filter(!is.na(mass), !is.na(height), !is.na(birth_year)) %>%
  summarise(min(mass), max(height), median(birth_year))
```

```
## # A tibble: 1 x 3
##   'min(mass)' 'max(height)' 'median(birth_year)'
##       <dbl>       <int>         <dbl>
## 1         17         228          46.5
```

## Question 10 (15 points)

In the flights data set, tailnum identifies the particular plane that was used in each flight. Throughout this question, use pipes.

- First, modify the flights data set so that it only contains flights that are not cancelled. Throughout this question, we'll only want to consider flights that were not cancelled; you can either store these non-cancelled flights in a new tibble, or do this same modification in each part below.

```
flights %>%
  filter(!is.na(dep_time), !is.na(arr_time))
```

```
## # A tibble: 328,063 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
## 7  2013     1     1     555             600          -5     913             854
## 8  2013     1     1     557             600          -3     709             723
## 9  2013     1     1     557             600          -3     838             846
## 10 2013     1     1     558             600          -2     753             745
## # ... with 328,053 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

- Group the flights data by tailnum. How many different tail numbers are there?

```
flights %>%
  filter(!is.na(dep_time), !is.na(arr_time)) %>%
  group_by(tailnum)
```

```
## # A tibble: 328,063 x 19
## # Groups:   tailnum [4,037]
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>        <dbl>    <int>         <int>
## 1  2013     1     1     517           515          2      830           819
## 2  2013     1     1     533           529          4      850           830
## 3  2013     1     1     542           540          2      923           850
## 4  2013     1     1     544           545         -1     1004          1022
## 5  2013     1     1     554           600         -6      812           837
## 6  2013     1     1     554           558         -4      740           728
## 7  2013     1     1     555           600         -5      913           854
## 8  2013     1     1     557           600         -3      709           723
## 9  2013     1     1     557           600         -3      838           846
## 10 2013     1     1     558           600         -2      753           745
## # ... with 328,053 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

4037 groups

- c. Make a summary tibble that shows, for each tailnum, how many times that plane flew out of new york and what the average departure delay of those flights was.

```
flights %>%
  filter(!is.na(dep_time), !is.na(arr_time)) %>%
  filter(origin == "JFK" | origin == "LGA" | origin == "EWR") %>%
  group_by(tailnum) %>%
  summarise(n_rows = n(), avg_dep = mean(dep_delay))
```

```
## # A tibble: 4,037 x 3
##   tailnum n_rows avg_dep
##   <chr>    <int>   <dbl>
## 1 D942DN      4  31.5
## 2 NOEGMQ    354   8.49
## 3 N10156    146  17.8
## 4 N102UW     48    8
## 5 N103US     46 -3.20
## 6 N104UW     46  10.1
## 7 N10575    271  22.3
## 8 N105UW     45   2.58
## 9 N107US     41 -0.463
## 10 N108UW     60   4.22
## # ... with 4,027 more rows
```

- d. What tailnum made the most flights out of new york in 2013?

```

flights %>%
  filter(!is.na(dep_time), !is.na(arr_time)) %>%
  filter(year == 2013)%>% filter( origin == "JFK" | origin == "LGA" | origin == "EWR") %>% group_by(tailnum)

## # A tibble: 1 x 3
##   tailnum n_rows avg_dep
##   <chr>    <int>   <dbl>
## 1 N725MQ      546    6.87

```

- e. Filter your tibble from (a) so that it only contains tail numbers that made at least 100 flights out of New York.

```

flights %>%
  filter(!is.na(dep_time), !is.na(arr_time)) %>%
  filter(origin == "JFK" | origin == "LGA" | origin == "EWR") %>%
  group_by(tailnum) %>% filter(n() >= 100)

## # A tibble: 222,870 x 19
## # Groups:   tailnum [1,210]
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>    <int>         <int>
## 1  2013     1     1     517           515           2      830           819
## 2  2013     1     1     533           529           4      850           830
## 3  2013     1     1     544           545          -1     1004          1022
## 4  2013     1     1     554           558          -4      740           728
## 5  2013     1     1     555           600          -5      913           854
## 6  2013     1     1     557           600          -3      709           723
## 7  2013     1     1     557           600          -3      838           846
## 8  2013     1     1     558           600          -2      849           851
## 9  2013     1     1     558           600          -2      853           856
## 10 2013     1     1     558           600          -2      923           937
## # ... with 222,860 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>

```


## Question 11 (25 points)

- a. (3 points) Make a bar chart showing how many flights in this data set departed from each of the three origin airports. (Hint: the easiest way to do this does not involve any data transformations)

```


flights %>%
  filter(!is.na(dep_time), !is.na(arr_time)) %>%
  filter(origin == "JFK" | origin == "LGA" | origin == "EWR") %>%
  ggplot(aes(x = origin)) + geom_bar()

```




- b. (6 points) Make a bar chart showing how many total minutes of departure delay there were for each of the three origin airports.

```
flights %>% filter(!is.na(dep_delay)) %>%  
  filter(origin == "JFK" | origin == "LGA" | origin == "EWR") %>%  
  group_by(origin) %>% summarize(total_delay = sum(dep_delay)) %>%  
  ggplot(mapping = aes(x = origin, y = total_delay)) + geom_col()
```



- c. (8 points) For each origin airport, compute the average delay of all flights leaving that airport. Show these average delays in a bar chart. Be sure you are removing any NA values *before* computing your averages.

```
flights %>%  
  filter(!is.na(dep_delay)) %>%  
  filter(origin == "JFK" | origin == "LGA" | origin == "EWR") %>%  
  group_by(origin) %>% summarize(average_delay = mean(dep_delay)) %>%  
  ggplot(mapping = aes(x = origin, y = average_delay)) + geom_col()
```



- d. (6 points) For each origin airport, compute the total number of flights, the number of flights that are cancelled, and the fraction of flights that are cancelled. A useful command is `sum(is.na(dep_time))`, which counts how many flights have a `dep_time` that is NA, that is, how many flights were cancelled. (You can visualize these in a bar chart if you want but it's not required)

```
flights %>%
  filter(origin == "JFK" | origin == "LGA" | origin == "EWR") %>%
  group_by(origin) %>%
  summarize(number_Cancelled_flights = sum(is.na(dep_time)),
            number_good_flights = sum(!is.na(dep_time)),
            number_flights = n(),
            Percent_cancelled = 100*(number_Cancelled_flights/number_flights))
```

```
## # A tibble: 3 x 5
##   origin number_Cancelled_fl~ number_good_fligh~ number_flights Percent_cancell~
##   <chr>          <int>          <int>          <int>          <dbl>
## 1 EWR             3239          117596         120835           2.68
## 2 JFK             1863          109416         111279           1.67
## 3 LGA             3153          101509         104662           3.01
```

- e. (2 points) Based on your answers to the previous parts, if you want to minimize your flight delays, which new york city area airport should you fly out of? If you want to minimize the chance that your flight is cancelled, what airport should you fly out of?



To minimize delays one should use the laguardia airport, and to minimize likelihood of cancellation, one should use John f. Kennedy airport.